

Beginner NumPy Practice Problems – Solutions

1. Create 1D and 2D arrays

Solution:

```
import numpy as np
arr1 = np.array([2, 4, 6, 8]) # 1D array
arr2 = np.array([[1, 2, 3],
[4, 5, 6]]) # 2x3 array
arr1.shape -> (4,)
arr2.shape -> (2, 3)
```

2. Check basic properties

Solution:

```
X = np.array([[1, 2, 3],
[4, 5, 6]])
print(X.shape) # (2, 3)
print(X.ndim) # 2 (2D array)
print(X.dtype) # usually int64 or int32 depending on system
```

3. Zeros and ones

Solution:

```
zeros_3x3 = np.zeros((3, 3))
ones_2x4 = np.ones((2, 4))
```

4. arange vs linspace

Solution:

```
a = np.arange(0, 10, 2) # [0 2 4 6 8]
b = np.linspace(0, 1, 5) # [0. 0.25 0.5 0.75 1.]
```

5. Indexing single elements

Solution:

```
X = np.array([[800, 2, 15],
[950, 2, 10],
[1200, 3, 8],
[1500, 3, 5]])
first_element = X[0, 0] # 800
third_row_bedrooms = X[2, 1] # 3
last_element = X[-1, -1] # 5
```

6. Row and column selection

Solution:

```
first_row = X[0, :] # [800, 2, 15]
second_column = X[:, 1] # [2, 2, 3, 3]
```

7. Slicing rows and columns

Solution:

```
first_two_rows = X[:2, :] # rows with 800 and 950
rows_1_to_3_cols_0_to_1 = X[1:4, 0:2]
```

```
# [[ 950, 2],  
# [1200, 3],  
# [1500, 3]]
```

8. Feature/target split

Solution:

```
X = data[:, :3] # features: size, bedrooms, age  
y = data[:, 3] # target: price  
X.shape -> (4, 3), y.shape -> (4,)
```

9. Mean and standard deviation

Solution:

```
overall_mean = data.mean()  
overall_std = data.std()  
Numerical values (approx):  
overall_mean ≈ 27781.12  
overall_std ≈ 49878.97
```

10. Mean per column

Solution:

```
col_means = data.mean(axis=0)  
Column means (approx):  
size mean ≈ 1112.50  
bedrooms mean ≈ 2.50  
age mean ≈ 9.50  
price mean ≈ 110000.00
```

The mean of the last column is the average house price.

11. Min and max per feature

Solution:

```
feature_mins = data[:, :3].min(axis=0)  
feature_maxs = data[:, :3].max(axis=0)  
Minimums (size, bedrooms, age): [800. 2. 5.]  
Maximums (size, bedrooms, age): [1500. 3. 15.]
```

12. Reshape and flatten

Solution:

```
a = np.array([1, 2, 3, 4, 5, 6])  
a_2x3 = a.reshape(2, 3) # [[1 2 3], [4 5 6]]  
a_flat = a_2x3.flatten() # [1 2 3 4 5 6]
```

13. Vertical and horizontal stacking

Solution:

```
A = np.array([[1, 2],  
[3, 4]])  
B = np.array([[5, 6],  
[7, 8]])  
v_stacked = np.vstack((A, B))  
# [[1 2],
```

```
# [3 4],  
# [5 6],  
# [7 8]]  
h_stacked = np.hstack((A, B))  
# [[1 2 5 6],  
# [3 4 7 8]]
```

14. Scalar operations

Solution:

```
X = np.array([10, 20, 30, 40])  
X1 = X - 10 # [0, 10, 20, 30]  
X2 = X / 10 # [1., 2., 3., 4.]
```

15. Elementwise operations

Solution:

```
a = np.array([1, 2, 3])  
b = np.array([10, 20, 30])  
a_plus_b = a + b # [11, 22, 33]  
a_times_b = a * b # [10, 40, 90]  
a_div_b = a / b # [0.1, 0.1, 0.1]
```

These are called elementwise operations because NumPy applies the operation on each pair of elements at the same index.

16. Vector dot product

Solution:

```
w = np.array([1, 2, 3])  
x = np.array([4, 5, 6])  
dot1 = np.dot(w, x)  
dot2 = w @ x  
# dot1 = dot2 = 1*4 + 2*5 + 3*6 = 4 + 10 + 18 = 32
```

17. Matrix–vector multiplication (simple model)

Solution:

```
X = np.array([[1, 2, 3],  
[4, 5, 6]])  
w = np.array([0.1, 0.2, 0.3])  
b = 0.5  
y_pred = X @ w + b  
# First row: 1*0.1 + 2*0.2 + 3*0.3 + 0.5 = 1.9  
# Second row: 4*0.1 + 5*0.2 + 6*0.3 + 0.5 = 3.7  
So y_pred = [1.9, 3.7]. Each element is a predicted output for one row (one data point).
```

18. Transpose and sums

Solution:

```
A = np.array([[1, 2, 3],  
[4, 5, 6]])  
A_T = A.T  
# A_T = [[1, 4],  
# [2, 5],
```

```
# [3, 6]
sum_axis0 = A.sum(axis=0) # [5, 7, 9]
sum_axis1 = A.sum(axis=1) # [6, 15]
axis=0: sum down each column.
axis=1: sum across each row.
```

19. Boolean indexing (filtering)

Solution:

```
prices = np.array([75000, 90000, 120000, 155000, 60000])
mask = prices > 100000
# mask -> [False, False, True, True, False]
expensive = prices[mask] # [120000, 155000]
count_expensive = expensive.size # 2
```

20. Shuffling dataset (X and y together)

Solution:

```
X = data[:, :3]
y = data[:, 3]
indices = np.random.permutation(len(X))
X_shuffled = X[indices]
y_shuffled = y[indices]
```

Now the rows in X_shuffled and the labels in y_shuffled are still matched but in random order.